

Infrastructure Document for Real Time Streaming

1. Overview

This Docker Compose deployment sets up a complete Apache Kafka ecosystem with the following components:

- **ZooKeeper**: Coordination service for Kafka brokers.
- **Kafka Broker**: Message broker storing and serving topics.
- **Schema Registry**: Manages Avro schemas for Kafka messages.
- **Kafka Connect**: Integration framework for streaming data between Kafka and external systems.
- **Control Center**: Confluent monitoring and management UI (default **UI**).
- **Kafka UI**: Third-party web interface to browse Kafka topics and manage clusters(Optional).
- **PostgreSQL & MySQL**: Source databases for Debezium connectors and Destination Database.
- **Spark**: JupyterLab environment for data processing and analytics.

All services communicate over a dedicated Docker bridge network (kafka-network). Persistent data is stored on the host via mounted volumes.

3. Services

3.1 Zookeeper

```
zookeeper:
  image: confluentinc/cp-zookeeper:6.2.15
  hostname: zookeeper
  container_name: zookeeper
  ports:
    - "2181:2181"
  volumes:
    - ./Data/zookeeper/data:/var/lib/zookeeper/data
    - ./Data/zookeeper/log:/var/lib/zookeeper/log
  environment:
    ZOOKEEPER_CLIENT_PORT: 2181
    ZOOKEEPER_TICK_TIME: 2000
    ZOOKEEPER_DATA_DIR: /var/lib/zookeeper/data
    ZOOKEEPER_DATA_LOG_DIR: /var/lib/zookeeper/log
  networks:
    - kafka-network
```

3.2 Kafka Broker

```
broker:
  image: confluentinc/cp-server:6.2.15
  hostname: broker
  container_name: broker
  depends_on:
    - zookeeper
  ports:
    - "9092:9092"
    - "29092:29092"
    - "9101:9101"
```

```

volumes:
  - ./connectors:/connectors # source kafka connect files
  - ./Data/kafka-data:/var/lib/kafka/data # Mount a host directory for Kafka data persistence

environment:
  KAFKA_BROKER_ID: 1
  KAFKA_ZOOKEEPER_CONNECT: 'zookeeper:2181'
  KAFKA_ADVERTISED_LISTENERS: PLAINTEXT_INTERNAL://broker:29092,PLAINTEXT://localhost:9092
  KAFKA_LISTENER_SECURITY_PROTOCOL_MAP: PLAINTEXT:PLAINTEXT,PLAINTEXT_INTERNAL:PLAINTEXT
  KAFKA_LISTENERS: PLAINTEXT_INTERNAL://0.0.0.0:29092,PLAINTEXT://0.0.0.0:9092
  KAFKA_INTER_BROKER_LISTENER_NAME: PLAINTEXT_INTERNAL
  KAFKA_METRIC_REPORTERS: io.confluent.metrics.reporter.ConfluentMetricsReporter
  KAFKA_OFFSETS_TOPIC_REPLICATION_FACTOR: 1
  KAFKA_GROUP_INITIAL_REBALANCE_DELAY_MS: 0
  KAFKA_CONFLUENT_LICENSE_TOPIC_REPLICATION_FACTOR: 1
  KAFKA_CONFLUENT_BALANCER_TOPIC_REPLICATION_FACTOR: 1
  KAFKA_TRANSACTION_STATE_LOG_MIN_ISR: 1
  KAFKA_TRANSACTION_STATE_LOG_REPLICATION_FACTOR: 1
  KAFKA_JMX_PORT: 9101
  KAFKA_JMX_HOSTNAME: localhost
  KAFKA_CONFLUENT_SCHEMA_REGISTRY_URL: http://schema-registry:8081
  CONFLUENT_METRICS_REPORTER_BOOTSTRAP_SERVERS: broker:29092
  CONFLUENT_METRICS_REPORTER_TOPIC_REPLICAS: 1
  CONFLUENT_METRICS_ENABLE: 'true'
  CONFLUENT_SUPPORT_CUSTOMER_ID: 'anonymous'
  KAFKA_AUTO_CREATE_TOPICS_ENABLE: 'true'
  KAFKA_DEFAULT_REPLICATION_FACTOR: 1
networks:
  - kafka-network

```

3.2 Schema Registry

```

schema-registry:
  image: confluentinc/cp-schema-registry:6.2.15
  hostname: schema-registry
  container_name: schema-registry
  depends_on:
    - broker
  ports:
    - "8081:8081"
  environment:
    SCHEMA_REGISTRY_HOST_NAME: schema-registry
    SCHEMA_REGISTRY_KAFKASTORE_BOOTSTRAP_SERVERS: 'broker:29092'
    SCHEMA_REGISTRY_LISTENERS: http://0.0.0.0:8081
  networks:
    - kafka-network

```

3.3 Kafka Connect

```

connect:
  image: confluentinc/cp-kafka-connect:6.2.15
  hostname: connect
  container_name: connect
  depends_on:
    - broker
    - schema-registry
  ports:

```

```

- "8083:8083"
volumes:
- ./connectors:/home/appuser/connectors # source kafka connect files
- ./plugins/debezium-connector/debezium-connector-mysql:/usr/share/java/debezium-connector-mysql
- ./plugins/debezium-connector/debezium-connector-postgres:/usr/share/java/debezium-connector-postgres
- ./plugins/kafka-connect-jdbc:/usr/share/java/kafka-connect-jdbc
environment:
CONNECT_BOOTSTRAP_SERVERS: 'broker:29092'
CONNECT_REST_ADVERTISED_HOST_NAME: connect
CONNECT_REST_PORT: 8083
CONNECT_GROUP_ID: compose-connect-group
CONNECT_CONFIG_STORAGE_TOPIC: docker-connect-configs
CONNECT_CONFIG_STORAGE_REPLICATION_FACTOR: 1
CONNECT_OFFSET_FLUSH_INTERVAL_MS: 1000
CONNECT_OFFSET_STORAGE_TOPIC: docker-connect-offsets
CONNECT_OFFSET_STORAGE_REPLICATION_FACTOR: 1
CONNECT_STATUS_STORAGE_TOPIC: docker-connect-status
CONNECT_STATUS_STORAGE_REPLICATION_FACTOR: 1
CONNECT_KEY_CONVERTER: org.apache.kafka.connect.storage.StringConverter
CONNECT_VALUE_CONVERTER: io.confluent.connect.avro.AvroConverter
CONNECT_VALUE_CONVERTER_SCHEMA_REGISTRY_URL: http://schema-registry:8081
CLASSPATH: /usr/share/java/kafka-connect-jdbc/*
CONNECT_PRODUCER_INTERCEPTOR_CLASSES: "io.confluent.monitoring.clients.interceptor.MonitoringProducerInterceptor"
CONNECT_CONSUMER_INTERCEPTOR_CLASSES: "io.confluent.monitoring.clients.interceptor.MonitoringConsumerInterceptor"
CONNECT_PLUGIN_PATH: "/usr/share/java,/usr/share/java/kafka-connect-jdbc,/usr/share/confluent-hub-components"
CONNECT_LOG4J_LOGGERS: org.apache.zookeeper=ERROR,org.I0Itec.zkclient=ERROR,org.reflections=ERROR
networks:
- kafka-network

```

3.4 Confluent Control Center(UI)

```

control-center:
image: confluentinc/cp-enterprise-control-center:6.2.15
hostname: control-center
container_name: control-center
depends_on:
- broker
- schema-registry
- connect
ports:
- "9021:9021"
environment:
CONTROL_CENTER_CONNECT_CONNECT-DEFAULT_CLUSTER_NAME: "Dev Connect Cluster"
CONTROL_CENTER_BOOTSTRAP_SERVERS: 'broker:29092'
CONTROL_CENTER_CONNECT_CONNECT-DEFAULT_CLUSTER: 'http://connect:8083'
CONTROL_CENTER_CONNECT_CLUSTER: 'http://connect:8083'
CONTROL_CENTER_SCHEMA_REGISTRY_URL: "http://schema-registry:8081"
CONTROL_CENTER_REPLICATION_FACTOR: 1
CONTROL_CENTER_INTERNAL_TOPICS_PARTITIONS: 1
CONTROL_CENTER_MONITORING_INTERCEPTOR_TOPIC_PARTITIONS: 1
CONFLUENT_METRICS_TOPIC_REPLICATION: 1
PORT: 9021
networks:
- kafka-network

```

3.5 Kafka UI (optional)

```
kafka-ui:
  image: provectuslabs/kafka-ui:latest
  container_name: kafka-ui
  depends_on:
    - broker
  environment:
    KAFKA_CLUSTERS_0_NAME: local
    KAFKA_CLUSTERS_0_BOOTSTRAPSERVERS: 'broker:29092'
    KAFKA_CLUSTERS_0_ZOOKEEPER: 'zookeeper:2181'
  ports:
    - "8080:8080"
  networks:
    - kafka-network
```

3.6 PostgreSQL

```
postgres:
  image: postgres:15
  container_name: postgres
  ports:
    - "5432:5432"
  environment:
    POSTGRES_USER: admin
    POSTGRES_PASSWORD: admin
    POSTGRES_DB: test_db
  volumes:
    - ./Data/postgres:/var/lib/postgresql/data
    - ./sql/init.sql:/docker-entrypoint-initdb.d/init.sql
  command:
    - "postgres"
    - "-c"
    - "wal_level=logical"
    - "-c"
    - "max_replication_slots=10"
    - "-c"
    - "max_wal_senders=10"
  healthcheck:
    test: ["CMD", "pg_isready", "-U", "demo_user"]
    interval: 10s
    timeout: 5s
    retries: 5
  networks:
    - kafka-network
```

3.7 MYSQL

```
mysql:
  image: mysql:8.0
  container_name: mysql
  ports:
    - "23306:3306"
  environment:
    MYSQL_ROOT_PASSWORD: root
    MYSQL_DATABASE: test_db
    MYSQL_USER: admin
    MYSQL_PASSWORD: admin
```

```

volumes:
  - ./Data/mysql:/var/lib/mysql
healthcheck:
  test: ["CMD", "mysqladmin", "ping", "-h", "localhost"]
  interval: 10s
  timeout: 5s
  retries: 5
networks:
  - kafka-network

```

3.8 Spark (Jupyter Notebook)

```

spark:
  build:
  context: ./docker/spark
  container_name: spark
  ports:
    - "8888:8888" # JupyterLab
    - "4040:4040" # Spark Web UI
  volumes:
    - ./Data/spark/notebooks:/home/jovyan/work
    - ./Data/spark/other_data/ivy2-cache:/home/jovyan/.ivy2
    - ./Data/spark/other_data/m2-repo:/home/jovyan/.m2/repository
    - ./plugins/kafka-connect-jdbc/mysql-connector-java-8.0.33.jar:/home/jovyan/work/jars/mysql-connector-java-8.0.33.jar
    - ./plugins/kafka-connect-jdbc/postgresql-42.7.7.jar:/home/jovyan/work/jars/postgresql-42.7.7.jar
  environment:
    - SPARK_DRIVER_MEMORY=2g
    - SPARK_EXECUTOR_MEMORY=2g
    - JUPYTER_TOKEN=spark@456
  command: >
    start-notebook.sh
    --NotebookApp.token='spark@456'
    --NotebookApp.ip='0.0.0.0'
    --NotebookApp.port=8888
    --NotebookApp.notebook_dir='/home/jovyan/work'
  networks:
    - kafka-network

```

Dockerfile

```

FROM jupyter/pyspark-notebook:latest
COPY requirements.txt /tmp/requirements.txt
RUN pip install --no-cache-dir -r /tmp/requirements.txt \
&& rm -f /tmp/requirements.txt || true

```

Requirement.txt

```

mysql-connector-python==8.3.0
Faker
ipywidgets
jupyterlab-execute-time
pyspark-stubs

```

```
fastavro
requests
streamlit
streamlit-autorefresh
```

4. Volumes and Persistent Data

- Host directories under `./Data/` for zookeeper, Kafka, MySQL, PostgreSQL, and spark caches
 - All the required plugin files are under `./plugins`
 - `debezium-connector-postgres`
 - `kafka-connect-jdbc-10.3.7.jar`
 - `mysql-connector-java-8.0.33.jar`
 - `postgresql-42.7.7.jar`
-

5. Scaling and Production Considerations

- **Replication:** Use replication factor > 1 for high availability
 - **ZK Replacement:** Consider Kafka KRaft mode
 - **Monitoring:** Integrate with Prometheus/Grafana
 - **Backup:** Snapshot volumes regularly
-