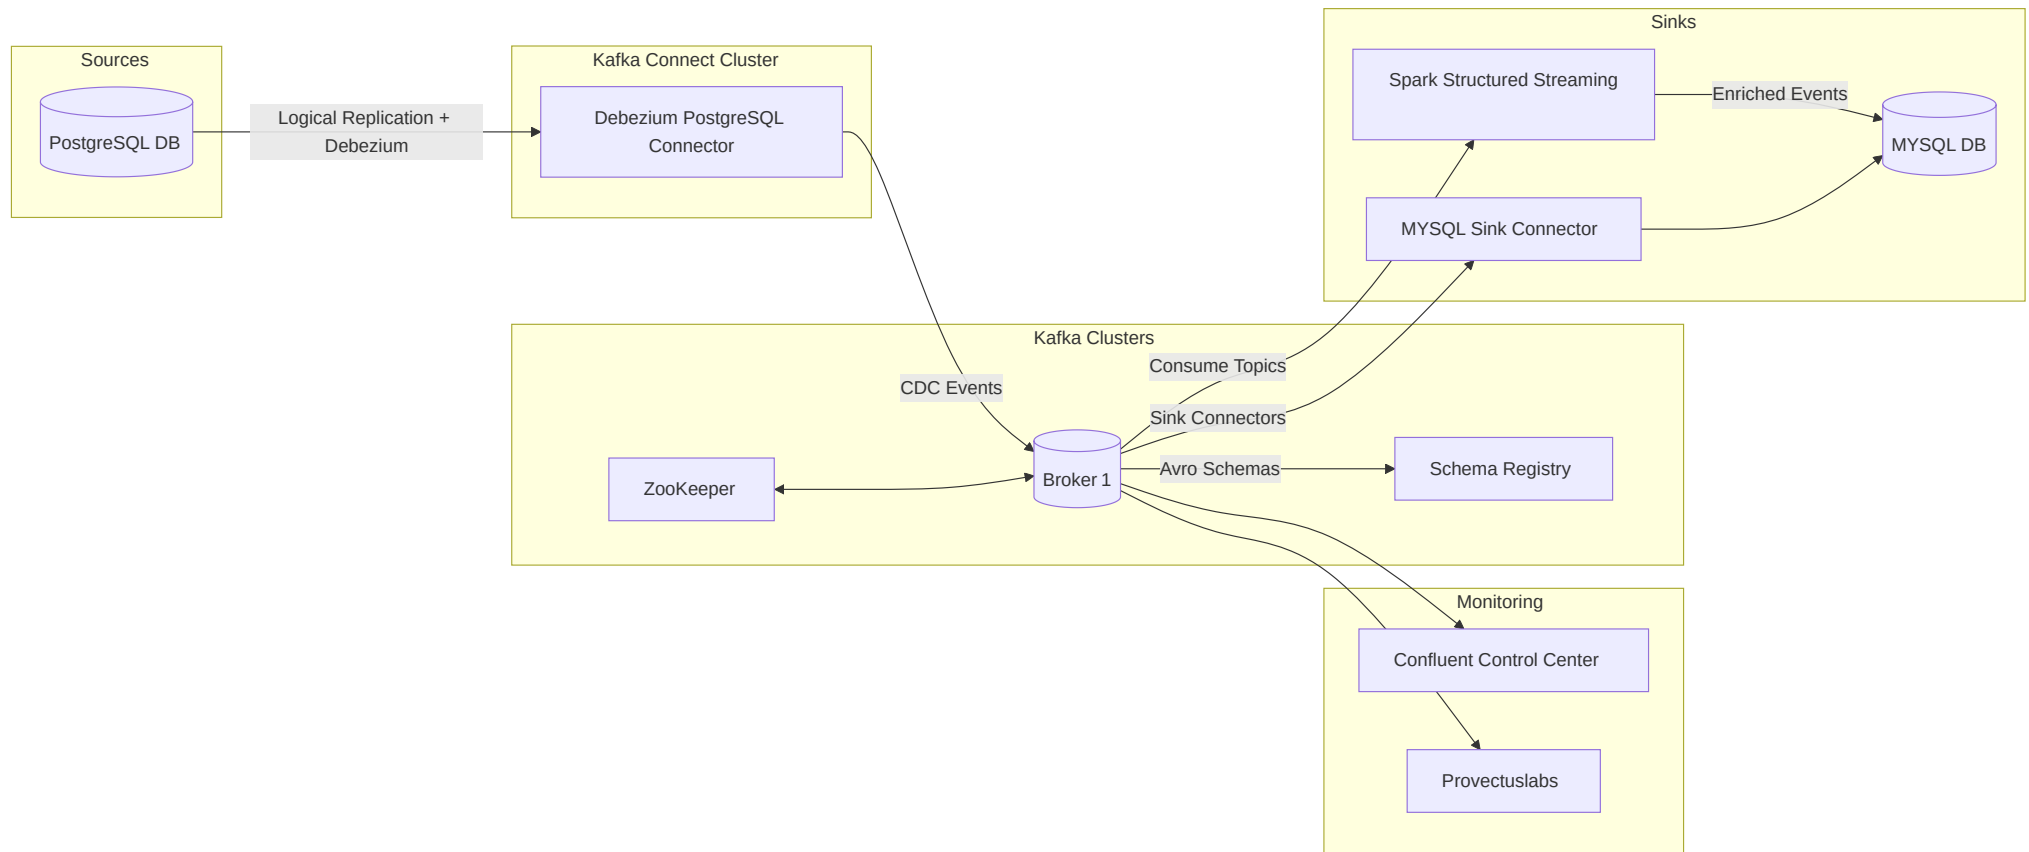


Real Time Streaming

Architecture Diagram of Real Time Streaming



1. Introduction

This document explains how a real-time streaming pipeline ingests change events from source databases, transports them through Kafka, processes them in-flight, and delivers them to downstream systems. I'll cover each component, data flow, and key concepts that are needed to build and operate such a system in production environment.

2. High-Level Data Flow

- Change Data Capture (CDC)**
 - Source:** PostgreSQL
 - Connector:** Debezium PostgreSQL Source Connector
 - Mechanism:** Logical replication stream (WAL) → Debezium pulls and serializes each INSERT/UPDATE/DELETE into Kafka-friendly events.

2. Message Broker

- **ZooKeeper**: Coordinates the Kafka brokers, maintains metadata, and handles leader election.
- **Kafka Broker (Broker 1)**: Receives CDC events into the `orders` topic, persists them across partitions and replicas, and serves them to consumers.

3. Schema Management

- **Schema Registry**: Stores Avro (or JSON/Protobuf) schemas for all topics. Producers register schemas on write; consumers fetch schemas on read to ensure compatibility.

4. Stream Processing

- **Spark Structured Streaming**: Consumes one or more Kafka topics, applies transformations/enrichments, and writes results directly to a sink.

5. Sink Connectors

- **MySQL Sink Connector**: Subscribes to Kafka topics and writes events into a downstream MySQL database, maintaining up-to-date tables.

6. Monitoring & Management

- **Confluent Control Center and Provectuslabs Kafka-UI**: Tracks cluster health, topic throughput, consumer lags, and connector status.

3. Component Breakdown

3.1 Change Data Capture (CDC)

- **Debezium PostgreSQL connector**: PostgreSQL uses `Write-Ahead Log (WAL)` to record every change before applying it to the actual data. We set `wal = logical` so that PostgreSQL allows row-level changes (insert/update/delete) to be decoded into logical events that CDC tools can consume. Without logical, PostgreSQL would only replicate raw binary file changes (physical replication), which cannot be easily parsed into meaningful per-row CDC events. Although **WAL** always logs all changes, logical decoding will only emit changes for tables which are defined in **publication**, Others are filtered out.

- **Logical Decoding Plugin:**

Decodes WAL records into a logical format, **pgoutput**: Native plugin used by PostgreSQL for logical replication

- **Replication Slot:**

Acts like a bookmark for a replication client. WAL logs needed by the slot will not be deleted until consumed — ensures no data loss. Different replication slot for different source connector.

```
SELECT * FROM pg_create_logical_replication_slot('debezium_slot', 'pgoutput');
```

- **Publication:**

Specifies **which tables** (and changes) are exposed via logical replication.

```
ALTER ROLE admin WITH REPLICATION;  
CREATE PUBLICATION my_publication FOR ALL TABLES;
```

```
- **PostgreSQL Debezium parameters (postgres_transactions.json)**
```

```
```json  
"bootstrap.servers": "broker:9092", # Kafka broker hostname and port
"connector.class": "io.debezium.connector.postgresql.PostgresConnector",
"database.hostname": "postgres",
"database.port": "5432",
"database.user": "admin",
"database.password": "admin",
"database.dbname": "transactions",
"max.batch.size": "10", # Maximum number of events to read
"max.queue.size": "20",
"plugin.name": "pgoutput", # Default plugin
"publication.name": "my_publication", # Publication name defined earlier
"slot.name": "transactions", # Replication Slot name
```

```
"table.include.list": "public.transaction_details", # List of table to be loaded
```

```
"database.server.name": "dev",
"topic.prefix": "dev",
"offset.storage.topic": "dev_connect_offsets",
"config.storage.topic": "dev_connect_configs",
"status.storage.topic": "dev_connect_status",
"offset.storage.replication.factor": "1",
"config.storage.replication.factor": "1",
"status.storage.replication.factor": "1",

"database.history.kafka.topic": "dev_schema_history",
"database.history.kafka.bootstrap.servers": "broker:9092",
"database.history.kafka.replication.factor": "1",

"key.converter": "io.confluent.connect.avro.AvroConverter",
"key.converter.schema.registry.url": "http://schema-registry:8081",
"value.converter": "io.confluent.connect.avro.AvroConverter",
"value.converter.schema.registry.url": "http://schema-registry:8081",
"snapshot.mode": "initial",
"time.precision.mode": "connect",

"transforms": "unwrap",
"transforms.unwrap.type": "io.debezium.transforms.ExtractNewRecordState",
"transforms.unwrap.drop.tombstones": "false",
"transforms.unwrap.delete.handling.mode": "rewrite",
"transforms.unwrap.add.fields": "op,source_table,ts_ms",
"include.schema.changes": "false"
```

```
PostgreSQL source connector file is present at : `/connectors/source/postgres_transactions.json`
```

To run a Kafka Connect source connector, like Debezium for PostgreSQL, we use a POST API request to the Kafka Connect REST endpoint.

```
curl -X POST -H "Content-Type: application/json" http://localhost:8083/connectors -d @postgres_transactions.json
```

To check the status of Kafka connect:

```
curl -s http://localhost:8083/connectors/connector_name/status
```

---

## 3.2 Kafka Broker & ZooKeeper

- **Broker:**
  - Hosts one or more **topics** divided into **partitions** for parallelism.
  - Each partition has a single *leader* (handles all reads/writes) and zero or more *followers* (replicate leader data).
  - *Replication factor*  $\geq 2$  ensures high availability.
- **ZooKeeper:**
  - Manages broker membership, topic metadata, and handles leader elections.

- In production, run an odd-numbered ensemble (3 or 5 nodes) for quorum.
- **Key parameters:**
  - *ISR (In-Sync Replicas)*: Followers that have fully caught up to the leader.
  - *min.insync.replicas*: Minimum followers that must acknowledge a write to consider it successful.

### 3.3 Schema Registry

- **Role:**
    - Centralizes schema definitions for Avro/Protobuf/JSON.
    - Validates producer schemas against compatibility rules (BACKWARD, FORWARD, FULL).
    - Allows consumers to evolve without code changes by retrieving schema by subject and version.
  - **Key Concepts:**
    - *Schema Compatibility*: Prevents breaking changes in production.
    - *Subject Naming*: `<topic>-key` and `<topic>-value`.
- 

### 3.4 Stream Processing (Spark Structured Streaming)

- **Overview**

This Spark application is designed to process financial transaction data ingested from a Kafka topic(`dev.public.transaction_details`). It performs the following tasks:
- Ingest Avro-encoded Kafka records
- Deserialize messages using Confluent Schema Registry
- Filter and aggregate data in 15-minute windows
- Write aggregated data to a MySQL database using an upsert strategy

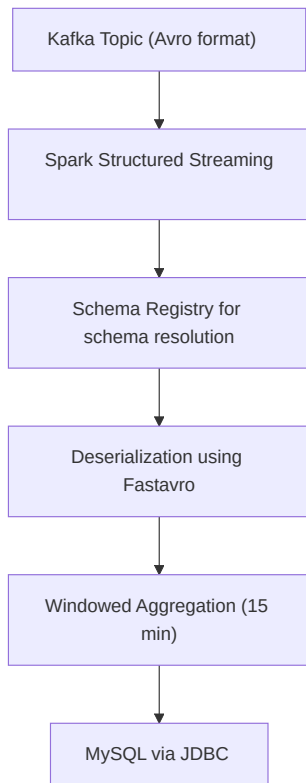


Fig: Flow diagram of spark consumer

- **Window Aggregation:** Uses a 5-minute watermark to handle out-of-order events, Aggregates over **15-minute** sliding windows

```
agg_df = (
 df.withWatermark("last_modified_date", "5 minutes")
 .groupBy(F.window(col("last_modified_date"), "15 minutes"))
 .agg(
 F.sum("amount").alias("total_amount"),
 F.count("amount").alias("txn_count")
)
)
```

- **Writing to MySQL with Upserts :** `foreachBatch` writes to MySQL using the following strategy:
  - Append mode to a staging table (`rp_stage_agg_15min_transactions`)
  - Upsert into final table (`rp_agg_15min_transactions`) using `ON DUPLICATE KEY UPDATE`
  - Truncate staging table after each batch

```
INSERT INTO rp_agg_15min_transactions (window_start, window_end, txn_count, total_amount)
SELECT window_start, window_end, txn_count, total_amount FROM rp_stage_agg_15min_transactions
ON DUPLICATE KEY UPDATE
window_end = VALUES(window_end),
txn_count = VALUES(txn_count),
```

```
total_amount = VALUES(total_amount);

TRUNCATE TABLE rp_stage_agg_15min_transactions;
```

- **Streaming Trigger and Checkpointing:**
  - Triggers every 1 minute for low-latency aggregation
  - Checkpointing ensures fault tolerance
  - Output mode is update, meaning only changed rows are emitted

```
.trigger(processingTime="1 minute")
.option("checkpointLocation", "checkpoint/gaurav/agg_15_min_n")
.outputMode("update")
```

All the results processed by spark structured streaming are dumped into rp database.  
Intermediate Stage in **MYSQL**

```
SELECT * FROM rp rp_stage_agg_15min_transactions;
```

Final processed table in **MYSQL**

```
SELECT * FROM rp rp_agg_15min_transactions order by window_start desc;
```


window_start	window_end	total_amoun	txn_coun	created_at
2025-07-22 09:30:00	2025-07-22 09:45:00	90495.77	16	2025-07-22 09:43:10
2025-07-22 09:15:00	2025-07-22 09:30:00	114473.29	22	2025-07-22 09:43:10
2025-07-22 09:00:00	2025-07-22 09:15:00	164494.95	29	2025-07-22 09:43:10
2025-07-22 08:45:00	2025-07-22 09:00:00	109653.91	24	2025-07-22 09:43:10
2025-07-22 08:30:00	2025-07-22 08:45:00	30375.25	9	2025-07-22 09:43:10
2025-07-22 04:45:00	2025-07-22 05:00:00	113291.77	23	2025-07-22 04:58:34
2025-07-22 04:30:00	2025-07-22 04:45:00	179630.80	33	2025-07-22 04:55:28
2025-07-22 04:15:00	2025-07-22 04:30:00	122291.23	26	2025-07-22 04:55:28
2025-07-22 04:00:00	2025-07-22 04:15:00	100418.04	18	2025-07-22 04:55:28

Streamlit UI is available after running file available at  
/home/jovyan/work/gaurav/streamlit/real\_time\_transactions.py

```
Streamlit run real_time_transactions.py
```

# Real Time Streaming

## 15 minutes Aggregation of Transaction Amount



	window_start	window_end	total_amount	txn_count	created_at
0	2025-07-22 09:30:00	2025-07-22 09:45:00	90495.77	16	2025-07-22 09:43:10
1	2025-07-22 09:15:00	2025-07-22 09:30:00	114473.29	22	2025-07-22 09:43:10
2	2025-07-22 09:00:00	2025-07-22 09:15:00	164494.95	29	2025-07-22 09:43:10
3	2025-07-22 08:45:00	2025-07-22 09:00:00	109653.91	24	2025-07-22 09:43:10
4	2025-07-22 08:30:00	2025-07-22 08:45:00	30375.25	9	2025-07-22 09:43:10
5	2025-07-22 04:45:00	2025-07-22 05:00:00	113291.77	23	2025-07-22 04:58:34
6	2025-07-22 04:30:00	2025-07-22 04:45:00	179630.8	33	2025-07-22 04:55:28

### 3.5 Kafka Connect Sink

- **MySQL Sink Connector:**
  - Pulls from Kafka topics, converts message format, and writes to MySQL tables.
  - Can perform upserts, deletes, and configurable batching.
- **Key Concepts:**
  - *Offset Storage:* Connect workers commit their progress back to a Kafka internal topic, allowing exactly-once delivery.

• **MYSQL Sink parameters**  
MySQL sink connector file is present at : /connectors/sink/mysql\_transactions.json

```
"connector.class": "io.confluent.connect.jdbc.JdbcSinkConnector",
"tasks.max": "1",
"topics": "dev.public.transaction_details",

"connection.url": "jdbc:mysql://mysql:3306/transactions?
useSSL=false&serverTimezone=UTC&allowPublicKeyRetrieval=true&sessionVariables=sql_mode='ALLOW_INVALID_DATES,ERROR_FOR_DIVISION_BY_ZERO,NO_ENGINE_SUBSTITUTION'",
"connection.user": "root",
"connection.password": "root",
```

```

"insert.mode": "upsert",
"auto.create": "true",
"auto.evolve": "true",
"pk.mode": "record_key",
"pk.fields": "txn_id",
"delete.enabled": "true",
"batch.size": "3000",
"delete.handling.mode": "rewrite",

"key.converter": "io.confluent.connect.avro.AvroConverter",
"key.converter.schema.registry.url": "http://schema-registry:8081",
"value.converter": "io.confluent.connect.avro.AvroConverter",
"value.converter.schema.registry.url": "http://schema-registry:8081",

"transforms": "unwrap,route",
"transforms.unwrap.type": "io.debezium.transforms.ExtractNewRecordState",
"transforms.unwrap.drop.tombstones": "true",

"transforms.route.type": "org.apache.kafka.connect.transforms.RegexRouter",
"transforms.route.regex": "dev\\.public\\.(.*)",
"transforms.route.replacement": "$1"

```

To run a Kafka Connect sink connector, like JDBC for MySQL, we use a POST API request to the Kafka Connect REST endpoint.

```
curl -X POST -H "Content-Type: application/json" http://localhost:8083/connectors -d @mysql_transactions.json
```

real time data streaming into **MYSQL**

```
SELECT * FROM transactions.transaction_details;
```

txn_id	sender_account_i	receiver_account_i	amount	last_modified_date	product_ic	product_type_i	transactor_module_i	module_ic	status	ts_ms	__op	__deleted	test
100629703	886794	805711	2085.18	2025-07-15 08:11:20	116	13	7	1	3	1752570320191	c	false	0
100816087	510910	109517	7572.26	2025-07-15 08:49:41	524	2	3	4	2	1752569806500	c	false	0
102170866	806184	537095	5895.74	2025-07-22 03:42:29	777	19	1	3	3	1753160088485	c	false	0
103002651	343459	479371	5042.18	2025-07-15 08:38:37	616	19	7	2	2	1752569713522	c	false	0
103464565	345299	410346	3831.45	2025-07-15 08:31:13	518	12	5	1	2	1752569713522	c	false	0
103647205	677223	973721	1564.00	2025-07-15 08:11:45	136	7	7	2	1	1752568935961	c	false	0
104754670	685833	737735	7205.78	2025-07-16 03:42:05	796	17	1	4	2	1752638871260	c	false	0
105000122	886227	390698	4717.22	2025-07-15 08:19:10	879	14	6	1	2	1752569018065	c	false	0
105102199	652601	400035	581.71	2025-07-22 04:32:30	693	1	7	4	3	1753160315113	c	false	0
105985329	785742	142630	3312.00	2025-07-22 09:19:16	302	4	9	4	4	1753177289377	c	false	0
106612954	806722	115970	4297.75	2025-07-22 03:45:42	532	13	1	3	3	1753160088597	c	false	0

### 3.6 Monitoring & Management

- **Confluent Control Center and Provectuslabs kafka-ui:**
  - Visualizes broker topic throughput, and consumer lags.



- Monitors connector health and data drift.

UI for Apache Kafka 83b5a60 4/10/2024, 15:26:25

Dashboard

local

Brokers

Topics

Consumers

Topics / dev.public.transaction\_details

Produce Message

Overview Messages Consumers Settings Statistics

Seek Type Offset Offset Partitions All items are selected. Key Serde String Value Serde String Clear all Submit Oldest First

Search Add Filters

DONE 287 ms 14 KB 500 messages consumed

	Offset	Partition	Timestamp	Key Preview	Value Preview
+	966	0	8/5/2025, 23:10:17	=====	=====tr...
+	967	0	8/5/2025, 23:10:17	=====	
+	968	0	8/5/2025, 23:10:17	=====	=====true...
+	969	0	8/5/2025, 23:10:17	=====	
+	970	0	8/5/2025, 23:10:17	=====	=====true...

### 3.7 Run and Validation

- Run /connectors/source/postgres\_transactions.json , PostgreSQL debezium source connector file to start real time change data capture (CDC)

```
curl -X POST -H "Content-Type: application/json" http://localhost:8083/connectors -d @postgres_transactions.json
```

- Run /connectors/sink/mysql\_transactions.json , MYSQL connector file to ingest data into destination database(mysql, transactions db)

```
curl -X POST -H "Content-Type: application/json" http://localhost:8083/connectors -d @mysql_transactions.json
```

- Run /home/jovyan/work/gaurav/streamlit/data\_producer\_transactions.py file to ingest data into source database table

```
spark-submit --master local[2] --jars /home/jovyan/work/jars/mysql-connector-java-8.0.33.jar,/home/jovyan/work/jars/postgresql-42.7.7.jar data_producer_transaction.py
```

- Run /home/jovyan/work/gaurav/streamlit/real\_time\_transactions.py file to monitor real time update in streamlit dashboard

```
streamlit run real_time_transactions.py
```

- Run /home/jovyan/work/gaurav/streamlit/spark\_consumer.py to perform real time streaming transformation using spark

```
spark-submit --master local[*] --packages org.apache.spark:spark-sql-kafka-0-10_2.12:3.5.0,org.apache.spark:spark-avro_2.12:3.5.0 --jars /home/jovyan/work/jars/mysql-connector-java-8.0.33.jar,/home/jovyan/work/jars/postgresql-42.7.7.jar spark_consumer.py
```

## 4. Summary

This real-time streaming pipeline captures database changes in near-real time, reliably transports them through a resilient Kafka cluster, processes or enriches the data on-the-fly, and delivers it to downstream systems(mysql) with end-to-end monitoring. By adhering to the concepts and best practices above, we can build a scalable, fault-tolerant, and maintainable streaming architecture.